

JUDGMENT AT MACHINE SPEED

Cultivating critical thinking and human judgment in the AI-augmented enterprise.



Two in three employees who use AI at work rely on its output without evaluating accuracy.

CONTENTS

01

Executive summary

02

Adoption has outrun scrutiny

03

How outputs are handled

04

The judgment operating model

05

Roadmap: visibility before control

06

Five priorities and scorecard

EXECUTIVE SUMMARY

CONTENT IS CHEAP. JUDGMENT IS SCARCE.

Generative AI has moved the constraint on enterprise value. The question is no longer how fast teams can produce, but how well they can evaluate what they produce.

Most organizations deployed AI faster than they built the habits, controls and culture to scrutinize it. The result is a hidden risk position: errors embedded in deliverables, sensitive data flowing into unmanaged tools, and a workforce that distrusts AI yet depends on it.



58%

OF EMPLOYEES USE AI AT WORK REGULARLY


The goal is not frictionless automation. It is the right friction in the right places: machine speed where errors are cheap, human scrutiny where they are not.

02 · THE SITUATION


ADOPTION HAS OUTRUN SCRUTINY

AI is now embedded in core work, but the verification habits needed to use it safely are not. The gap between use and scrutiny is where enterprise risk accumulates.

The benefits are real. Among employees who use AI:



Higher efficiency



Better access to
information




Better work quality

The problem is not use. It is how outputs are handled once they arrive.

Source: KPMG & University of Melbourne, Trust, attitudes and use of AI: A global study 2025 (47 countries).


HOW OUTPUTS ARE HANDLED

Two in three rely on AI output without checking it



ADOPTION BY ECONOMY

Emerging economies use more, train more and worry less



Advanced economies are more sceptical but less trained: caution without capability, which tends to produce informal, ungoverned use rather than less use.

04 · THE JUDGMENT OPERATING MODEL

BUILD SCEPTICISM INTO THE WORKFLOW

Scepticism must be built into how work flows, not left to individual temperament. Match the level of review to the cost of an error.

3

RISK TIERS, EACH WITH ITS OWN
MINIMUM HUMAN REVIEW

TIER THE WORK BY RISK

1

LOW

**Drafts and
brainstorming**

Self-check; AI use noted.

2

MODERATE

**Internal decisions and
reports**

Named owner verifies
sources and key figures.

3

HIGH

**Client, financial, legal,
safety, strategy**

Assumption map plus
independent red-team
review before release.

RED-TEAM REVIEW

**Task-level checks**

Does the output meet the standard an expert would expect?

**Adversarial checks**

Probe for prompt injection and unverified citations.

**Component-level checks**

Are sources real, figures traceable, reasoning unbiased?


**No metric shortcut**

Automated scores can rate confidently wrong output highly.

MAP THE ASSUMPTIONS


Test the critical claims with the weakest evidence first

Break a fluent AI recommendation into testable claims: do customers want it, does it pay, can we deliver it, does it survive change? Then score each on importance and evidence.



Score evidence on what supports a claim, not how it reads. Fluent prose and model confidence count as weak evidence.

FOUR CHALLENGE PROMPTS



Caution: a model critiquing its own output shares its blind spots. Challenge prompts widen the search; they do not replace checks against primary sources.

05 · ROADMAP

VISIBILITY BEFORE CONTROL



Three phases, each ending at a gate leadership must pass before scaling. Visibility comes first, because controls designed blind push AI use further out of sight.

DAYS 0–30

Illuminate

Launch a 30-day AI amnesty · Map tools, data flows and demand · Define risk tiers 1–3 · Name an executive sponsor

◆ **Gate 1:** shadow-use map complete, tiers agreed

DAYS 30–90

Equip

License secure AI tools · Deploy access and data-loss controls · Train on bias and verification · Pilot red-team review in 2–3 high-risk workflows

◆ **Gate 2:** pilots catch errors without slowing delivery

DAYS 90–180

Embed

Scale controls to all tier-3 work · Add psychological safety to manager goals · Set time-reinvestment targets · Review KPIs quarterly

D180

What good looks like: AI use is visible, high-risk outputs are reviewed by a named owner, the KPI scorecard is baselined, and managers are measured on whether people feel safe to challenge AI.

06 · WHAT LEADERS SHOULD DO

FIVE PRIORITIES

- 01

Run a 30-day AI amnesty, then license secure alternatives

0–30 DAYS
- 02

Tier AI-assisted work by risk and set minimum review per tier

0–60 DAYS
- 03

Embed assumption mapping and red teams in high-stakes decisions

30–90 DAYS
- 04

Train for bias awareness and verification, not just prompting

30–120 DAYS
- 05

Hold managers to account for psychological safety and reinvesting saved time

60–180 DAYS

KPI SCORECARD

Visibility Tools disclosed ↑ · unsanctioned tools ↓	Verification Tier-3 outputs with documented review → 100%
Quality AI errors caught before release vs. after ↑	Security AI-related data-loss incidents and time to detect ↓
Psychological safety Comfort raising AI errors (team survey) ↑	Capability and value No-AI exercise scores stable · hours reinvested ↑

Set baselines in the first 30 days and review quarterly. Judgment quality is hard to measure directly, so track a basket of indicators rather than one metric.



MACHINE SPEED WHERE ERRORS ARE CHEAP. HUMAN SCRUTINY WHERE THEY ARE NOT.

The future belongs to organizations that treat AI as an environment demanding a higher calibre of human oversight, not as a replacement for human judgment.

KEY SOURCES

KPMG & University of Melbourne — Trust, attitudes and use of AI: A global study 2025

Deloitte — Workforce Trust with AI, TrustID Q3 2025

Reutlingen University — Mitigating Automation Bias in Generative AI Through Nudges

Harvard Business School — The Dangers of Deferring to AI

MIT Sloan — 'AI gravity' is pulling you toward dependency

Strategyzer — Assumptions mapping · Amy C. Edmondson — psychological safety

PREPARED BY ANUPAM BARUAH